

Appendix: Python Code in Jupyter Notebook

(Available at github.com/pi-prescott/spark-standalone)

In [1]: `# [Q1] first we initiate a SparkSession with the spark-master`

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("pyspark-notebook").\
    master("spark://spark-master:7077").\
    config("spark.executor.memory", "1024m").\
    getOrCreate()

spark
```

Out[1]: **SparkSession - in-memory**

SparkContext

[Spark UI](#)

Version	v3.0.0
Master	spark://spark-master:7077
AppName	pyspark-notebook

In [2]: `# [Q2a] load CSV into spark dataframe`

```
# and [Q2b] check schema
wrong_schema = spark.read.csv(path='../data/covid19.csv',header=True)
# by default the type of each column is apparently assumed to be String.
print(wrong_schema.printSchema())
```

```
root
|-- continent: string (nullable = true)
|-- location: string (nullable = true)
|-- date: string (nullable = true)
|-- total_cases: string (nullable = true)
|-- new_cases: string (nullable = true)
|-- total_deaths: string (nullable = true)
|-- new_deaths: string (nullable = true)
```

None

In [3]: `# if you ask explicitly, spark will try to infer the schema automatically`

```
infer_schema = spark.read.csv(
    path='../data/covid19.csv',header=True,inferSchema=True)
infer_schema.printSchema()
# in this case it gets the integers right, but just treats the date as a string
```

```
root
|-- continent: string (nullable = true)
|-- location: string (nullable = true)
|-- date: string (nullable = true)
|-- total_cases: integer (nullable = true)
|-- new_cases: integer (nullable = true)
|-- total_deaths: integer (nullable = true)
|-- new_deaths: integer (nullable = true)
```

In [4]: `# or you can specify the schema explicitly`

```
from pyspark.sql.types import (StructField,
                                StringType,
                                IntegerType,
                                DateType,
```

StructType)

```
data_schema = [StructField('continent',StringType(),True),
                StructField('location',StringType(),True),
                StructField('date',DateType(),True),
                StructField('total_cases',IntegerType(),True),
                StructField('new_cases',IntegerType(),True),
                StructField('total_deaths',IntegerType(),True),
                StructField('new_deaths',IntegerType(),True)]
```

```
correct_struct = StructType(fields=data_schema)
```

```
dataframe = spark.read.csv(
    path='../data/covid19.csv', header=True, schema=correct_struct)
```

```
# and we can confirm that this time the types are correct
print(dataframe.printSchema())
```

```
root
|-- continent: string (nullable = true)
|-- location: string (nullable = true)
|-- date: date (nullable = true)
|-- total_cases: integer (nullable = true)
|-- new_cases: integer (nullable = true)
|-- total_deaths: integer (nullable = true)
|-- new_deaths: integer (nullable = true)
```

None

In [5]: `# if we wanted to convert to the older-style RDD we easily could`

```
rdd = dataframe.rdd
print(f'Created `rdd` {type(rdd)} from `dataframe` {type(dataframe)}.')
# ... and vice versa
new_dataframe = rdd.toDF()
print(f'Created `new_dataframe` {type(new_dataframe)}'
      + f' from `rdd` {type(rdd)}.')
```

```
Created `rdd` <class 'pyspark.rdd.RDD'> from `dataframe` <class 'pyspark.sql.data
frame.DataFrame'>.
Created `new_dataframe` <class 'pyspark.sql.dataframe.DataFrame'> from `rdd` <cla
ss 'pyspark.rdd.RDD'>.
```

In [6]: `# the simplest way to drop null values from a spark 2.0 dataframe`

```
# ...is like this
drop_na = dataframe.dropna()
```

In [7]: `# [Q2c] but we can use the `.filter()` method if we like`

```
filtered_df = dataframe.filter(
    ' and '.join([f'{x} is not null' for x in dataframe.columns])
)
```

In [8]: `print(f'Before filtering we had {dataframe.count()} rows...')`

```
print(f'Using `.dropna()` leaves us {drop_na.count()} rows.')
print(f'Using `.filter()` leaves us {filtered_df.count()} rows.')
if drop_na.count() == filtered_df.count():
    print('Good, those numbers are the same!')
else:
    print('Not good -- those numbers should be the same...')
```

```
Before filtering we had 53087 rows...
Using `.dropna()` leaves us 39974 rows.
Using `.filter()` leaves us 39974 rows.
Good, those numbers are the same!
```

In [9]: `# [Q3] use aggregate and groupBy functions to see highest `total_deaths` in each`

```
hi_total_deaths = filtered_df.groupBy('location')\
    .agg({'total_deaths': 'max'})
```

```
In [10]: hi_total_deaths.show()
```

```
+-----+
|location|max(total_deaths)|
+-----+
|Chad|96|
|Paraguay|1327|
|Russia|26589|
|Yemen|600|
|Senegal|322|
|Sweden|5918|
|Guyana|119|
|Jersey|32|
|Philippines|7053|
|Djibouti|61|
|Malaysia|238|
|Singapore|28|
|Fiji|2|
|Turkey|9950|
|United States Vir...|21|
|Western Sahara|1|
|Malawi|183|
|Iraq|10724|
|Sint Maarten (Dut...|22|
|Germany|10183|
+-----+
only showing top 20 rows
```

```
In [11]: # the assignment suggests that the number of total_deaths for Sweden should be 986
# however it is actually 5918
hi_total_deaths.filter(hi_total_deaths.location=='Sweden').show()
```

```
+-----+
|location|max(total_deaths)|
+-----+
|Sweden|5918|
+-----+
```

```
In [12]: # however, we would get the result 986 if we hadn't
# explicitly made sure to load the CSV with the correct schema
wrong_schema.groupBy('location').agg({'total_deaths': 'max'})\
    .filter(wrong_schema.location=='Sweden').show()
```

```
+-----+
|location|max(total_deaths)|
+-----+
|Sweden|986|
+-----+
```

```
In [13]: # [Q4] use max and min functions to see which country
# has highest and lowest 'total_cases'
# NB: 'total_cases' are given for every date,
# so for country with lowest can't simply find min(total_cases)
# as we'll get an earlier date with a lower figure
# rather than the country with the lowest final total_cases
# -- however, this is obviously not an issue for the maximum figure
import pyspark.sql.functions as F

filtered_df.select(F.max('total_cases')).show()
filtered_df.groupBy('location').max('total_cases')\
    .select(F.min('max(total_cases)')).show()
```

```
+-----+
|max(total_cases)|
+-----+
|8779653|
+-----+
```

```
+-----+
|min(max(total_cases))|
+-----+
|13|
+-----+
```

```
In [14]: # to see a list of the countries with the highest and lowest total_cases count...
total_cases = filtered_df.groupBy('location').max('total_cases')
print('Countries with Highest Total Number of Cases')
total_cases.orderBy('max(total_cases)',ascending=False).show()
print('Countries with Lowest Total Number of Cases')
total_cases.orderBy('max(total_cases)',ascending=True).show()
```

Countries with Highest Total Number of Cases

```
+-----+
|location|max(total_cases)|
+-----+
|United States|8779653|
|India|7990322|
|Brazil|5439641|
|Russia|1547774|
|France|1198695|
|Spain|1116738|
|Argentina|1116596|
|Colombia|1033218|
|United Kingdom|917575|
|Mexico|901268|
|Peru|892497|
|South Africa|717851|
|Iran|581824|
|Italy|564778|
|Chile|504525|
|Germany|464239|
|Iraq|459908|
|Bangladesh|401586|
|Indonesia|396454|
|Philippines|373144|
+-----+
```

only showing top 20 rows

Countries with Lowest Total Number of Cases

```
+-----+
|location|max(total_cases)|
+-----+
|Montserrat|13|
|Fiji|33|
|British Virgin Is...|72|
|Northern Mariana ...|92|
|Antigua and Barbuda|124|
|Brunei|148|
|Bonaire Sint Eust...|150|
|Bermuda|194|
|Barbados|233|
|Cayman Islands|239|
|Guernsey|266|
|Monaco|320|
|Isle of Man|352|
|Mauritius|439|
|Liechtenstein|483|
|Tanzania|509|
|Comoros|517|
|Taiwan|550|
|Burundi|558|
|Jersey|560|
+-----+
```

only showing top 20 rows